

Boston house price prediction

The problem that we are going to solve here is that given a set of features that describe a house in Boston, our machine learning model must predict the house price. To train our machine learning model with boston housing data, we will be using scikit-learn's boston dataset.

In this dataset, each row describes a boston town or suburb. There are 506 rows and 13 attributes (features) with a target column (price). <https://archive.ics.uci.edu/ml/machine-learning-databases/housing/housing.names> (<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/housing.names>)

```
In [1]: 1 # Importing the libraries
        2 import pandas as pd
        3 import numpy as np
        4 from sklearn import metrics
        5 import matplotlib.pyplot as plt
        6 import seaborn as sns
        7 %matplotlib inline
```

```
In [2]: 1 # Importing the Boston Housing dataset
        2 from sklearn.datasets import load_boston
        3 boston = load_boston()
```

```
In [3]: 1 # Initializing the dataframe
        2 data = pd.DataFrame(boston.data)
```

```
In [4]: 1 # See head of the dataset
        2 data.head()
```

```
Out[4]:
```

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
|---|---------|------|------|-----|-------|-------|------|--------|-----|-------|------|--------|------|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | 396.90 | 4.98 |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | 396.90 | 9.14 |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | 392.83 | 4.03 |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.63 | 2.94 |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | 396.90 | 5.33 |

```
In [5]: 1 #Adding the feature names to the dataframe
        2 data.columns = boston.feature_names
        3 data.head()
```

Out[5]:

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT |
|---|---------|------|-------|------|-------|-------|------|--------|-----|-------|---------|--------|-------|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | 396.90 | 4.0 |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | 396.90 | 9.0 |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | 392.83 | 4.0 |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.63 | 2.0 |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | 396.90 | 5.0 |

CRIM per capita crime rate by town

ZN proportion of residential land zoned for lots over 25,000 sq.ft.

INDUS proportion of non-retail business acres per town

CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)

NOX nitric oxides concentration (parts per 10 million)

RM average number of rooms per dwelling

AGE proportion of owner-occupied units built prior to 1940

DIS weighted distances to five Boston employment centres

RAD index of accessibility to radial highways

TAX full-value property-tax rate per 10,000usd

PTRATIO pupil-teacher ratio by town

B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town

LSTAT % lower status of the population

Each record in the database describes a Boston suburb or town.

```
In [6]: 1 #Adding target variable to dataframe
        2 data['PRICE'] = boston.target
        3 # Median value of owner-occupied homes in $1000s
```

```
In [7]: 1 #Check the shape of dataframe
        2 data.shape
```

Out[7]: (506, 14)

```
In [8]: 1 data.columns
```

```
Out[8]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TA
X',
              'PTRATIO', 'B', 'LSTAT', 'PRICE'],
              dtype='object')
```

```
In [9]: 1 data.dtypes
```

```
Out[9]: CRIM      float64
        ZN        float64
        INDUS     float64
        CHAS      float64
        NOX       float64
        RM        float64
        AGE       float64
        DIS       float64
        RAD       float64
        TAX       float64
        PTRATIO   float64
        B         float64
        LSTAT     float64
        PRICE     float64
dtype: object
```

```
In [10]: 1 # Identifying the unique number of values in the dataset
        2 data.nunique()
```

```
Out[10]: CRIM      504
        ZN         26
        INDUS      76
        CHAS        2
        NOX        81
        RM        446
        AGE        356
        DIS        412
        RAD         9
        TAX        66
        PTRATIO    46
        B         357
        LSTAT     455
        PRICE     229
dtype: int64
```

```
In [11]: 1 # Check for missing values
        2 data.isnull().sum()
```

```
Out[11]: CRIM      0
        ZN         0
        INDUS      0
        CHAS       0
        NOX        0
        RM         0
        AGE        0
        DIS        0
        RAD        0
        TAX        0
        PTRATIO    0
        B          0
        LSTAT      0
        PRICE      0
dtype: int64
```

```
In [13]: 1 # See rows with missing values
        2 data[data.isnull().any(axis=1)]
```

Out[13]:

| CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT | PRICE |
|------|----|-------|------|-----|----|-----|-----|-----|-----|---------|---|-------|-------|
|------|----|-------|------|-----|----|-----|-----|-----|-----|---------|---|-------|-------|

```
In [14]: 1 # Viewing the data statistics
        2 data.describe()
```

Out[14]:

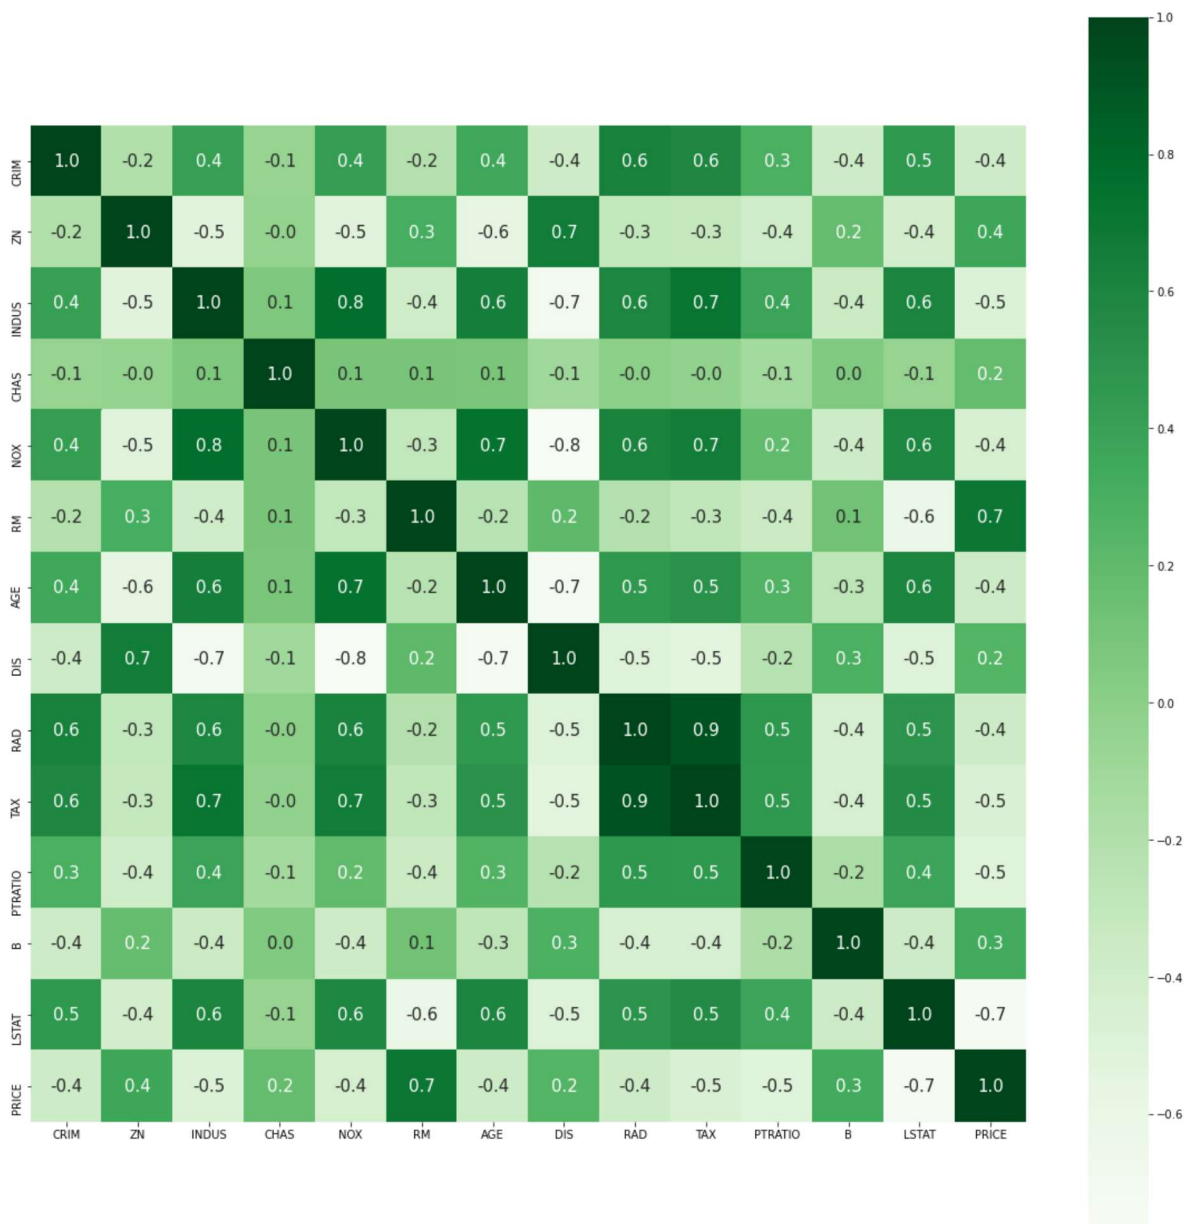
| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | |
|--------------|------------|------------|------------|------------|------------|------------|------------|-----|
| count | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506 |
| mean | 3.613524 | 11.363636 | 11.136779 | 0.069170 | 0.554695 | 6.284634 | 68.574901 | 3 |
| std | 8.601545 | 23.322453 | 6.860353 | 0.253994 | 0.115878 | 0.702617 | 28.148861 | 2 |
| min | 0.006320 | 0.000000 | 0.460000 | 0.000000 | 0.385000 | 3.561000 | 2.900000 | 1 |
| 25% | 0.082045 | 0.000000 | 5.190000 | 0.000000 | 0.449000 | 5.885500 | 45.025000 | 2 |
| 50% | 0.256510 | 0.000000 | 9.690000 | 0.000000 | 0.538000 | 6.208500 | 77.500000 | 3 |
| 75% | 3.677083 | 12.500000 | 18.100000 | 0.000000 | 0.624000 | 6.623500 | 94.075000 | 5 |
| max | 88.976200 | 100.000000 | 27.740000 | 1.000000 | 0.871000 | 8.780000 | 100.000000 | 12 |

```
In [15]: 1 # Finding out the correlation between the features
        2 corr = data.corr()
        3 corr.shape
```

Out[15]: (14, 14)

```
In [16]: 1 # Plotting the heatmap of correlation between features
2 plt.figure(figsize=(20,20))
3 sns.heatmap(corr, cbar=True, square=True, fmt='.1f', annot=True, annot_k
```

Out[16]: <AxesSubplot:>



```
In [17]: 1 # Splitting target variable and independent variables
2 X = data.drop(['PRICE'], axis = 1)
3 y = data['PRICE']
```

```
In [18]: 1 # Splitting to training and testing data
2
3 from sklearn.model_selection import train_test_split
4 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.3,
```

Linear regression

Training the model

```
In [19]: 1 # Import Library for Linear Regression
          2 from sklearn.linear_model import LinearRegression
          3
          4 # Create a Linear regressor
          5 lm = LinearRegression()
          6
          7 # Train the model using the training sets
          8 lm.fit(X_train, y_train)
```

Out[19]: LinearRegression()

```
In [20]: 1 # Value of y intercept
          2 lm.intercept_
```

Out[20]: 36.357041376595205

```
In [21]: 1 #Converting the coefficient values to a dataframe
          2 coefficients = pd.DataFrame([X_train.columns,lm.coef_]).T
          3 coefficients = coefficients.rename(columns={0: 'Attribute', 1: 'Coefficient'})
          4 coefficients
```

Out[21]:

| | Attribute | Coefficients |
|----|-----------|--------------|
| 0 | CRIM | -0.12257 |
| 1 | ZN | 0.055678 |
| 2 | INDUS | -0.008834 |
| 3 | CHAS | 4.693448 |
| 4 | NOX | -14.435783 |
| 5 | RM | 3.28008 |
| 6 | AGE | -0.003448 |
| 7 | DIS | -1.552144 |
| 8 | RAD | 0.32625 |
| 9 | TAX | -0.014067 |
| 10 | PTRATIO | -0.803275 |
| 11 | B | 0.009354 |
| 12 | LSTAT | -0.523478 |

Model Evaluation

```
In [22]: 1 # Model prediction on train data
          2 y_pred = lm.predict(X_train)
```

```
In [23]: 1 # Model Evaluation
          2 print('R^2:', metrics.r2_score(y_train, y_pred))
          3 print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_train, y_pred)) * (len(y_train) - 1) / (len(y_train) - 2))
          4 print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
          5 print('MSE:', metrics.mean_squared_error(y_train, y_pred))
          6 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

R²: 0.7465991966746854

Adjusted R²: 0.736910342429894

MAE: 3.08986109497113

MSE: 19.07368870346903

RMSE: 4.367343437774162

R² : It is a measure of the linear relationship between X and Y. It is interpreted as the proportion of the variance in the dependent variable that is predictable from the independent variable.


Adjusted R² :The adjusted R-squared compares the explanatory power of regression models that contain different numbers of predictors.

MAE : It is the mean of the absolute value of the errors. It measures the difference between two continuous variables, here actual and predicted values of y.


MSE: The mean square error (MSE) is just like the MAE, but squares the difference before summing them all instead of using the absolute value.

RMSE: The mean square error (MSE) is just like the MAE, but squares the difference before summing them all instead of using the absolute value.

```
In [24]: 1 # Visualizing the differences between actual prices and predicted values
2 plt.scatter(y_train, y_pred)
3 plt.xlabel("Prices")
4 plt.ylabel("Predicted prices")
5 plt.title("Prices vs Predicted prices")
6 plt.show()
```



```
In [25]: 1 # Checking residuals
2 plt.scatter(y_pred, y_train-y_pred)
3 plt.title("Predicted vs residuals")
4 plt.xlabel("Predicted")
5 plt.ylabel("Residuals")
6 plt.show()
```




There is no pattern visible in this plot and values are distributed equally around zero. So Linearity assumption is satisfied


```
In [26]: 1 # Checking Normality of errors
2 sns.distplot(y_train-y_pred)
3 plt.title("Histogram of Residuals")
4 plt.xlabel("Residuals")
5 plt.ylabel("Frequency")
6 plt.show()
```

C:\Users\admin\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)



Here the residuals are normally distributed. So normality assumption is satisfied

For test data

```
In [27]: 1 # Predicting Test data with the model
2 y_test_pred = lm.predict(X_test)
```

```
In [28]: 1 # Model Evaluation
2 acc_linreg = metrics.r2_score(y_test, y_test_pred)
3 print('R^2:', acc_linreg)
4 print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_test, y_test_pred))*(len(
5 print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
6 print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
7 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

R²: 0.7121818377409195

Adjusted R²: 0.6850685326005713

MAE: 3.8590055923707407

MSE: 30.053993307124127

RMSE: 5.482152251362974

Here the model evaluations scores are almost matching with that of train data. So the model is not overfitting.

Random Forest Regressor

Train the model

```
In [29]: 1 # Import Random Forest Regressor
          2 from sklearn.ensemble import RandomForestRegressor
          3
          4 # Create a Random Forest Regressor
          5 reg = RandomForestRegressor()
          6
          7 # Train the model using the training sets
          8 reg.fit(X_train, y_train)
```

Out[29]: RandomForestRegressor()

Model Evaluation

```
In [30]: 1 # Model prediction on train data
          2 y_pred = reg.predict(X_train)
```

```
In [31]: 1 # Model Evaluation
          2 print('R^2:', metrics.r2_score(y_train, y_pred))
          3 print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_train, y_pred)) * (len(y_train) - 1) / (len(y_train) - 2))
          4 print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
          5 print('MSE:', metrics.mean_squared_error(y_train, y_pred))
          6 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

R^2: 0.9785878017207631


Adjusted R^2: 0.9777691000218511

MAE: 0.8499576271186446


MSE: 1.611713929378534

RMSE: 1.2695329571848595

```
In [32]: 1 # Visualizing the differences between actual prices and predicted values
2 plt.scatter(y_train, y_pred)
3 plt.xlabel("Prices")
4 plt.ylabel("Predicted prices")
5 plt.title("Prices vs Predicted prices")
6 plt.show()
```



```
In [33]: 1 # Checking residuals
2 plt.scatter(y_pred, y_train-y_pred)
3 plt.title("Predicted vs residuals")
4 plt.xlabel("Predicted")
5 plt.ylabel("Residuals")
6 plt.show()
```



For test data

```
In [34]: 1 # Predicting Test data with the model
        2 y_test_pred = reg.predict(X_test)
```

```
In [35]: 1 # Model Evaluation
        2 acc_rf = metrics.r2_score(y_test, y_test_pred)
        3 print('R^2:', acc_rf)
        4 print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_test, y_test_pred))*(len(y_test)-1)/(len(y_test)-2))
        5 print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
        6 print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
        7 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

R^2: 0.8302539268272174

Adjusted R^2: 0.8142633547167379

MAE: 2.5325328947368426

MSE: 17.724897230263167

RMSE: 4.210094681864431

XGBoost Regressor

Training the model

```
In [36]: 1 # Import XGBoost Regressor
        2 from xgboost import XGBRegressor
        3
        4 #Create a XGBoost Regressor
        5 reg = XGBRegressor()
        6
        7 # Train the model using the training sets
        8 reg.fit(X_train, y_train)
```

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
<ipython-input-36-c5d0a9b9906d> in <module>
```

```
    1 # Import XGBoost Regressor
----> 2 from xgboost import XGBRegressor
      3
      4 #Create a XGBoost Regressor
      5 reg = XGBRegressor()
```

ModuleNotFoundError: No module named 'xgboost'

max_depth (int) – Maximum tree depth for base learners.

learning_rate (float) – Boosting learning rate (xgb's "eta")

n_estimators (int) – Number of boosted trees to fit.

gamma (float) – Minimum loss reduction required to make a further partition on a leaf node of the tree.

`min_child_weight` (int) – Minimum sum of instance weight(hessian) needed in a child.

`subsample` (float) – Subsample ratio of the training instance.

`colsample_bytree` (float) – Subsample ratio of columns when constructing each tree.

`objective` (string or callable) – Specify the learning task and the corresponding learning objective or a custom objective function to be used (see note below).

`nthread` (int) – Number of parallel threads used to run xgboost. (Deprecated, please use `n_jobs`)

`scale_pos_weight` (float) – Balancing of positive and negative weights

Model Evaluation

```
In [37]: 1 # Model prediction on train data
          2 y_pred = reg.predict(X_train)
```

```
In [38]: 1 # Model Evaluation
          2 print('R^2:',metrics.r2_score(y_train, y_pred))
          3 print('Adjusted R^2:',1 - (1-metrics.r2_score(y_train, y_pred))*(len(y_train)-1)/(len(y_train)-2))
          4 print('MAE:',metrics.mean_absolute_error(y_train, y_pred))
          5 print('MSE:',metrics.mean_squared_error(y_train, y_pred))
          6 print('RMSE:',np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

R^2: 0.9785878017207631


Adjusted R^2: 0.9777691000218511

MAE: 0.8499576271186446


MSE: 1.611713929378534

RMSE: 1.2695329571848595

```
In [39]: 1 # Visualizing the differences between actual prices and predicted values
2 plt.scatter(y_train, y_pred)
3 plt.xlabel("Prices")
4 plt.ylabel("Predicted prices")
5 plt.title("Prices vs Predicted prices")
6 plt.show()
```



```
In [40]: 1 # Checking residuals
2 plt.scatter(y_pred, y_train-y_pred)
3 plt.title("Predicted vs residuals")
4 plt.xlabel("Predicted")
5 plt.ylabel("Residuals")
6 plt.show()
```



For test data

```
In [41]: 1 #Predicting Test data with the model
        2 y_test_pred = reg.predict(X_test)
```

```
In [42]: 1 # Model Evaluation
        2 acc_xgb = metrics.r2_score(y_test, y_test_pred)
        3 print('R^2:', acc_xgb)
        4 print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_test, y_test_pred))*(len(y_test)-1)/(len(y_test)-2))
        5 print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
        6 print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
        7 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

R^2: 0.8302539268272174

Adjusted R^2: 0.8142633547167379

MAE: 2.5325328947368426

MSE: 17.724897230263167

RMSE: 4.210094681864431

SVM Regressor

```
In [43]: 1 # Creating scaled set to be used in model to improve our results
        2 from sklearn.preprocessing import StandardScaler
        3 sc = StandardScaler()
        4 X_train = sc.fit_transform(X_train)
        5 X_test = sc.transform(X_test)
```

Train the model

```
In [44]: 1 # Import SVM Regressor
        2 from sklearn import svm
        3
        4 # Create a SVM Regressor
        5 reg = svm.SVR()
```

```
In [45]: 1 # Train the model using the training sets
        2 reg.fit(X_train, y_train)
```

Out[45]: SVR()

C : float, optional (default=1.0): The penalty parameter of the error term. It controls the trade off between smooth decision boundary and classifying the training points correctly.

kernel : string, optional (default='rbf'): kernel parameters selects the type of hyperplane used to separate the data. It must be one of 'linear', 'poly', 'rbf', 'sigmoid', 'precomputed' or a callable.

degree : int, optional (default=3): Degree of the polynomial kernel function ('poly'). Ignored by all other kernels.

gamma : float, optional (default='auto'): It is for non linear hyperplanes. The higher the gamma value it tries to exactly fit the training data set. Current default is 'auto' which uses $1 / n_features$.

coef0 : float, optional (default=0.0): Independent term in kernel function. It is only significant in

Model Evaluation

```
In [46]: 1 # Model prediction on train data
         2 y_pred = reg.predict(X_train)
```

```
In [47]: 1 # Model Evaluation
         2 print('R^2:', metrics.r2_score(y_train, y_pred))
         3 print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_train, y_pred)) * (len(y_train) - 1) / (len(y_train) - 2))
         4 print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
         5 print('MSE:', metrics.mean_squared_error(y_train, y_pred))
         6 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

R^2: 0.6419097248941195


Adjusted R^2: 0.628218037904777

MAE: 2.9361501059460284


MSE: 26.953752101332935

RMSE: 5.191700309275655

```
In [48]: 1 # Visualizing the differences between actual prices and predicted values
         2 plt.scatter(y_train, y_pred)
         3 plt.xlabel("Prices")
         4 plt.ylabel("Predicted prices")
         5 plt.title("Prices vs Predicted prices")
         6 plt.show()
```




```
In [49]: 1 # Checking residuals
2 plt.scatter(y_pred,y_train-y_pred)
3 plt.title("Predicted vs residuals")
4 plt.xlabel("Predicted")
5 plt.ylabel("Residuals")
6 plt.show()
```



For test data

```
In [50]: 1 # Predicting Test data with the model
2 y_test_pred = reg.predict(X_test)
```

```
In [51]: 1 # Model Evaluation
2 acc_svm = metrics.r2_score(y_test, y_test_pred)
3 print('R^2:', acc_svm)
4 print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_test, y_test_pred))*(len(
5 print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
6 print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
7 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

R²: 0.5900158460478174

Adjusted R²: 0.5513941503856553

MAE: 3.7561453553021673

MSE: 42.81057499010247

RMSE: 6.542979060802691

Evaluation and comparison of all the models

```
In [52]: 1 models = pd.DataFrame({
2         'Model': ['Linear Regression', 'Random Forest', 'XGBoost', 'Support V
3         'R-squared Score': [acc_linreg*100, acc_rf*100, acc_xgb*100, acc_svm*
4         models.sort_values(by='R-squared Score', ascending=False)
```

Out[52]:

| | Model | R-squared Score |
|---|-------------------------|-----------------|
| 1 | Random Forest | 83.025393 |
| 2 | XGBoost | 83.025393 |
| 0 | Linear Regression | 71.218184 |
| 3 | Support Vector Machines | 59.001585 |